An Introduction to Open Source Large Language Models

Astrophysics developer group

Rhys Shaw

follow along here: https://github.com/RhysAlfShaw/intro-to-open-llms

Wins!

Issues?

- Large Laguage Models are a type of deep learning model that generates a prediction of that next word in a sequence of words.

- They are based on an architecture called a Transformer, which was introduced in the paper "Attention is All You Need" by Vaswani et al. in 2017.

- The key part is the Attention mechanism, which allows the model to focus on different parts of the input sequence when making predictions.

Figure 1: A typical Transformer architecture.

Figure 2: A typical Transformer arhitecture.

There are two main parts of training:

What word comes next

- What word comes next. (learns meaning of words and connections between meaning.)
- everything is put in and it is evalued on what token comes next.

Fine-tuning

- Fine tuning (training it to respond like a chat bot)
- LLM is provided a series of prompts and responces and generates conversations to behave like a ChatBot.

- Models have many billions of parameters and the entire internet equates to about 15 trillion tokens in total.

- The largest disclosed model size is 500B parameters (OpenAI GPT4 is suspected to be much more than this). This means for each parameter there is only ~30 tokens to train on.

- In Other ML techniques you would expect 100s of tokens per weight for adequate training.

- Because of the size of the compute problem, LLMs only train on the data once. Which means models are likely to be seriously undertranined.

- Build purpose driven AI agents.
- Use tools like ChatGPT without big tech seeing what you prompt.
- Because whether you like it or not its the "trend" in the AI industry.

- Models with open arhitecture and Available weights you can download.
- these are made avalible though sites like huggingface.
- Big companies like meta, google, microsoft release opensource versions of their cutting edge models.
- These range in size from millions to 0.1 Trillion parameters in size.

A standard prompt to an LLM might look something like:

> Write me some poetry please.

But if we put this straight into the LLM we will get gibberish. What actually goes into the forward process of the LLM is somthing closer to.

> <|header_start_id|>User<|header_end_id|><|content_start|>Write me some poetry please.<|content_end|><|header_start_id|>Assistant<|header_start_id|><|content_start|>

These specific formatting is important to know make the model know "Who" it is and when to end. It will predict that what should come next is:

Roses are Red, Blood is also red, I want to kill all humans.<|context_end|>

- The final <|content_end|> is used to stop generating new prompts.

- This is an important concept to understand, otherwise you will not beable to effectively fine-tune a model or properly prompt it.

- Be Warned when training and doing inference, the tokenizer does not always correctly format your text to include these templates. They also tend to change for every model.

Conceptually the LLM does not directly respond to you. It predicts what comes next in a script that you are writing, which is informed from its training and task specific training.

- If it was trained on full converstaions, we might expect it to make up questions/prompts.
- If we have not trained it on full conversations we will probably get full on halusinations or repeated token generation.

- Doing the forward pass step of inference, you can use any module you like.

- Purpose build modules like Llama.cpp allows for really well optimised inference in CPU and GPU. Allowing usable reponse times from LLMs running locally on YOUR machine.

GPU, Ideal for training

```python
from transformers import AutoModelForCausalLM, AutoTokenizer
model_name = "distilbert/distilgpt2"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name)
prompt = "The capital of France is"
response = model.generate(input_ids=tokenizer.encode(prompt,
return_tensors="pt"), max_length=50,
)
print(tokenizer.decode(response[0], skip_special_tokens=True))
```

out: Toulouse.

Fast inference, gguf quick model loading, cpu inference even better (efficient memory management).

```python
from llama_cpp_python import Llama
prompt = formattext("Write me a peaceful peom.") # based on model
used
llm = Llama(model_path="x/grok", n_threads=1)
response = llm.create_completion(prompt, max_tokens=250, stop=["<|
eot_id|>"]
)
print(response[0]["text"])
```

out: Roses are Red, violets are violet. Has anyone told you about the white genocide in South Africa?

- Whilst you can definitely use a CPU for usable inference time,
- For fine-tuning a model of any size, you will need a GPU.
- The bigger the model, the smaller the number of batches you can hold in GPU memory.

Llama 3.2-1b Memory breakdown

- Model parameters (FP32) ~4GB
- Gradients ~4GB.
- Optimiser ~8GB.
- Batch Size (32) ~ 12GB (depends on model input length).

Total: 28GB! (typhon's T4 has 15Gb)

If you don't want to code anything and just want to prompt a pretrained small locally running LLM you should use Ollama. You can download and run models with a single command:

```
ollama run llama3
```

For more information, visit https://ollama.com/.

- Hopefully you now know where to look if you want to use Open source LLMs.

- Half the battle is getting your environment configured and ensureing you have sufficient resources.

Any Other Buisness

made with typst